



Advanced HPC Part 1



Erin Shaw and Cesar Sul

Advanced Cyberinfrastructure Research and Education Facilitation

USC Center for High-Performance Computing

Calculate Pi

- Numerically evaluate pi
 - $\int_0^1 \frac{4}{\sqrt{1+x^2}} dx = \pi$
 - Sample size vs accuracy?
- Have to run many times



Calculate pi

- How to run

```
./calc_pi -n 100  
num_steps=100  
  
pi=3.1515759869231288
```

- Want to run from n=1 to n=1000
- Want to collect results into one file
- Want to run statistics on the one file

Job Arrays on Slurm

- Job Arrays let you reuse a job script
- Each job does a similar task
- Read more on [Slurm's Job Array documentation](#)
- Example files:
 - `/home/rcf-proj/workshop/adv_hpc/job_array`

Job Arrays on Slurm

```
#!/bin/bash

#SBATCH --ntasks=1
#SBATCH --cpus-per-task=1
#SBATCH --time=00:30:00
#SBATCH --export=none
#SBATCH --array=1-10

./pi_calc -n $SLURM_ARRAY_TASK_ID >
output${SLURM_ARRAY_TASK_ID}.txt
```

Job Arrays on Slurm

- Similar to normal job
- Include `#SBATCH --array=1-10`
- Job array elements differentiate themselves with `$SLURM_ARRAY_TASK_ID`

- [ttroj@hpc-login3 advanced_hpc]\$ myqueue

JOBID	USER	ACCOUNT	PARTITION	NAME	TASKS	CPUS_PER_TASK	MIN_ MEMORY	START_ TIME	TIME	TIME_ LIMIT	STATE	NODELIST (REASON)
3959285	ttroj	lc_ttl	scavenge	job_array.slurm	1	1	1G	0	0:00	30:00:00	CONFIGURING	hpc0682
3959286	ttroj	lc_ttl	scavenge	job_array.slurm	1	1	1G	0	0:00	30:00:00	CONFIGURING	hpc0682
3959287	ttroj	lc_ttl	scavenge	job_array.slurm	1	1	1G	0	0:00	30:00:00	CONFIGURING	hpc0682
3959288	ttroj	lc_ttl	scavenge	job_array.slurm	1	1	1G	0	0:00	30:00:00	CONFIGURING	hpc0682
3959289	ttroj	lc_ttl	scavenge	job_array.slurm	1	1	1G	0	0:00	30:00:00	CONFIGURING	hpc0682
3959290	ttroj	lc_ttl	scavenge	job_array.slurm	1	1	1G	0	0:00	30:00:00	CONFIGURING	hpc0682
3959291	ttroj	lc_ttl	scavenge	job_array.slurm	1	1	1G	0	0:00	30:00:00	CONFIGURING	hpc0682
3959292	ttroj	lc_ttl	scavenge	job_array.slurm	1	1	1G	0	0:00	30:00:00	CONFIGURING	hpc0682
3959293	ttroj	lc_ttl	scavenge	job_array.slurm	1	1	1G	0	0:00	30:00:00	CONFIGURING	hpc0683
3959284	ttroj	lc_ttl	scavenge	job_array.slurm	1	1	1G	0	0:00	30:00:00	CONFIGURING	hpc0683

Job Arrays on Slurm

- This works but it generates messy output
- Output files are organized alphabetically, not numerically

```
$grep pi output*.txt
```

```
output10.txt:pi=3.2399259889071588  
output1.txt:pi=4.000000000000000000  
output2.txt:pi=3.600000000000000001  
output3.txt:pi=3.4564102564102557  
output4.txt:pi=3.3811764705882354  
output5.txt:pi=3.3349261138109898  
output6.txt:pi=3.3036297331379298  
output7.txt:pi=3.2810484527641703  
output8.txt:pi=3.2639884944910889  
output9.txt:pi=3.2506461552653931
```


Job Arrays on Slurm

- We can "zero pad" our files
 - output1.txt -> output01.txt
 - Alphabetic sorting also becomes numeric sorting

- Change two lines in script

```
outfile=$(printf "output%02d.txt" $SLURM_ARRAY_TASK_ID)  
  
./pi_calc -n $SLURM_ARRAY_TASK_ID > $outfile
```

- %02d means take the variable `$SLURM_ARRAY_TASK_ID` and put as many zeros as required to make it 2 characters wide

Job Arrays on Slurm

```
#!/bin/bash

#SBATCH --ntasks=1
#SBATCH --cpus-per-task=1
#SBATCH --time=00:30:00
#SBATCH --export=none
#SBATCH --array=1-10

outfile=$(printf "output%02d.txt" $SLURM_ARRAY_TASK_ID)

./pi_calc -n $SLURM_ARRAY_TASK_ID > $outfile
```

Job Arrays on Slurm

- Organized files makes our lives easier down the line

```
$grep pi output*.txt
```

```
output01.txt:pi=4.000000000000000000  
output02.txt:pi=3.600000000000000001  
output03.txt:pi=3.4564102564102557  
output04.txt:pi=3.3811764705882354  
output05.txt:pi=3.3349261138109898  
output06.txt:pi=3.3036297331379298  
output07.txt:pi=3.2810484527641703  
output08.txt:pi=3.2639884944910889  
output09.txt:pi=3.2506461552653931  
output10.txt:pi=3.2399259889071588
```

Job Arrays on Slurm

- Organized files makes our lives easier down the line


```
$grep pi output*.txt | sed 's/^.*=/'
```

```
4.000000000000000000
3.600000000000000001
3.4564102564102557
3.3811764705882354
3.3349261138109898
3.3036297331379298
3.2810484527641703
3.2639884944910889
3.2506461552653931
3.2399259889071588
```

Look for lines
with 'pi' in
them. Display
only text after
'=' character

- We don't have to parse this later

Summary of results



Job Arrays on Slurm

■ Pros

- Only need 1 job script
- See all steps in 1 file
- (Can be) Easy to set up
- Many jobs can be queued up

■ Cons

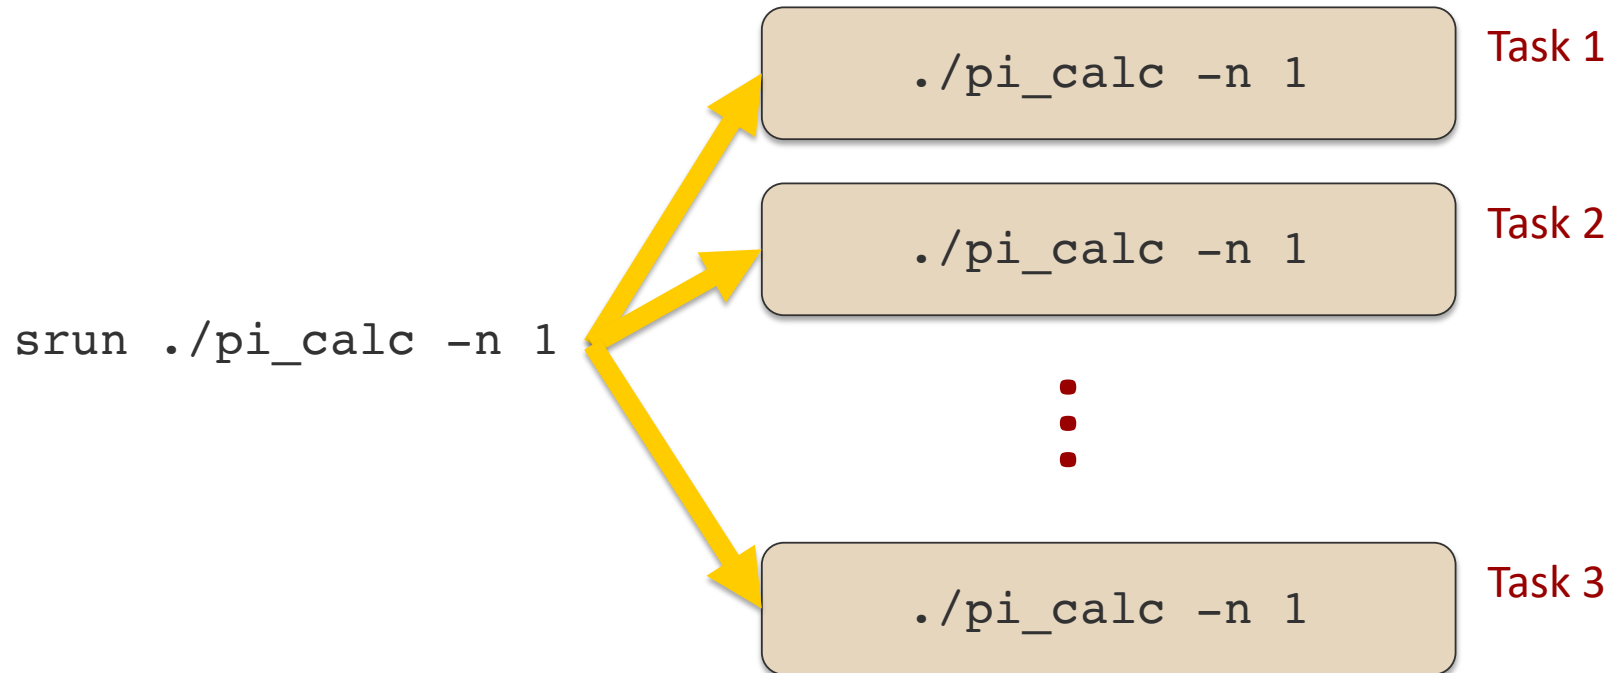
- 10 jobs running limit in main partition
- Node resources may sit idle (depending on program)

srun on Slurm

- Slurm's srun utility can launch parallel jobs
- `srun <command>` will launch `<command>` on all "tasks"

```
$ salloc --ntasks=4 --cpus-per-task=8
$ hostname
hpc0972
$ srun hostname
hpc0971
hpc0972
hpc0972
hpc0971
$ srun --ntasks=2 hostname
hpc0972
hpc0971
```

srun on Slurm



How to get unique behavior on each task?

srun on Slurm

```
#!/bin/bash

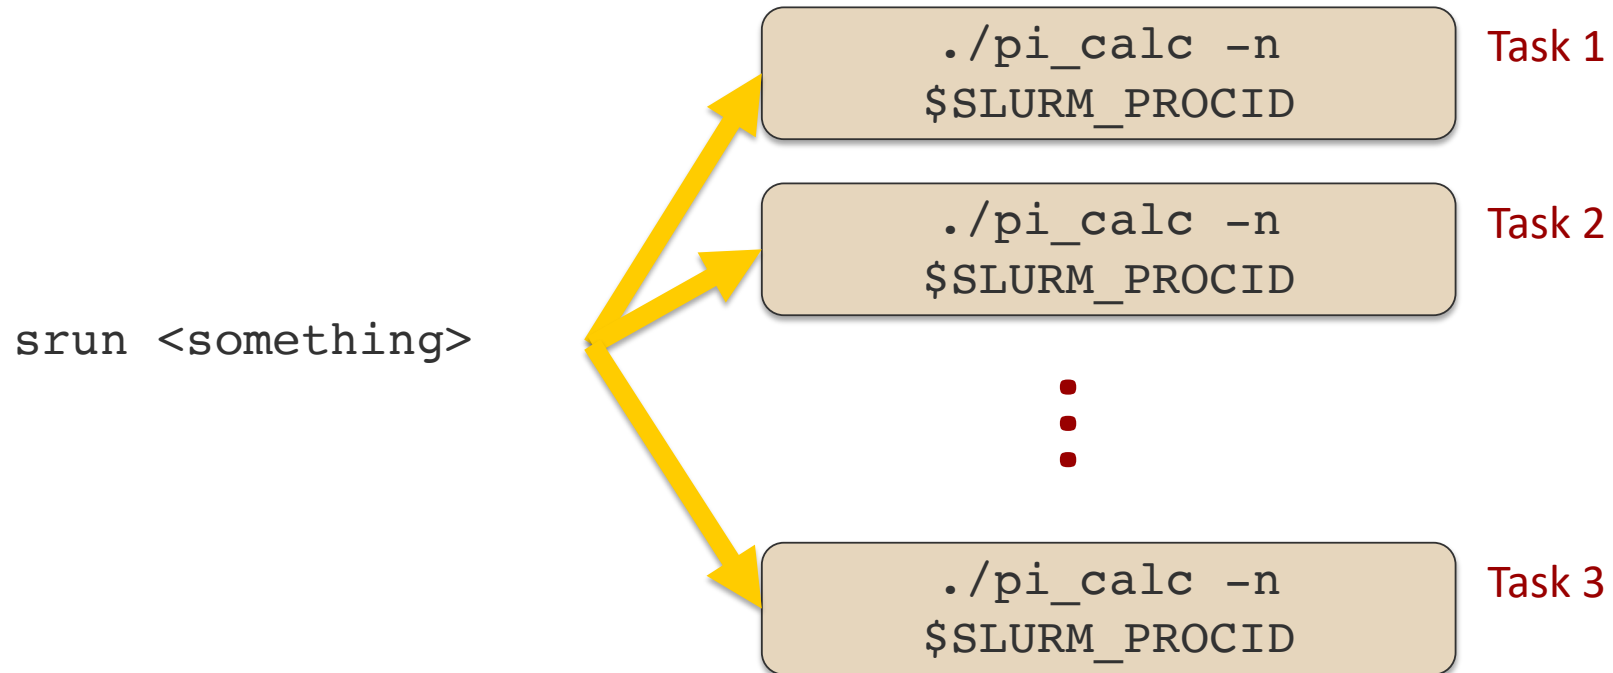
#SBATCH --ntasks=10
#SBATCH --cpus-per-task=1
#SBATCH --time=00:30:00
#SBATCH --export=none

srun ./pi_calc -n 1
```

srun and Tasks

- What is a "task"?
- According to Slurm it's a "process"
 - pi_calc is a process
- You can assign resources per process
 - cpus-per-task
 - mem-per-cpu
- Slurm will not split up tasks across nodes
- Just like \$SLURM_ARRAY_ID, tasks can have \$SLURM_PROCID
 - When launched with srun, \$SLURM_PROCID ranges from 0 to N-1

srun on Slurm



srun on Slurm

```
#!/bin/bash

#SBATCH --ntasks=10
#SBATCH --cpus-per-task=1
#SBATCH --time=00:30:00
#SBATCH --export=none

srun wrapper.sh
```

- We can't directly use `srun ./pi_calc -n $SLURM_PROCID`
- Outside of `srun`, `$SLURM_PROCID` is set to 0.
- Create "wrapper" script that determines unique work

srun on Slurm

```
#!/bin/bash
# All environment settings are initialized here

# stepsize of 0 is nonsensical, we must increment it
stepsize=$((SLURM_PROCID + 1))

# fancy zero padding like in the job array example
outfile=$(printf "output%02d.txt" $stepsize)

./pi_calc -n $stepsize > $outfile
```

- Each wrapper script has unique `$SLURM_PROCID`
- We had to process `$SLURM_PROCID` a little

Job Arrays on Slurm

- Organized files makes our lives easier down the line

```
$grep pi output*.txt
```

```
output01.txt:pi=4.000000000000000000  
output02.txt:pi=3.600000000000000001  
output03.txt:pi=3.4564102564102557  
output04.txt:pi=3.3811764705882354  
output05.txt:pi=3.3349261138109898  
output06.txt:pi=3.3036297331379298  
output07.txt:pi=3.2810484527641703  
output08.txt:pi=3.2639884944910889  
output09.txt:pi=3.2506461552653931  
output10.txt:pi=3.2399259889071588
```

srun on Slurm

■ Pros

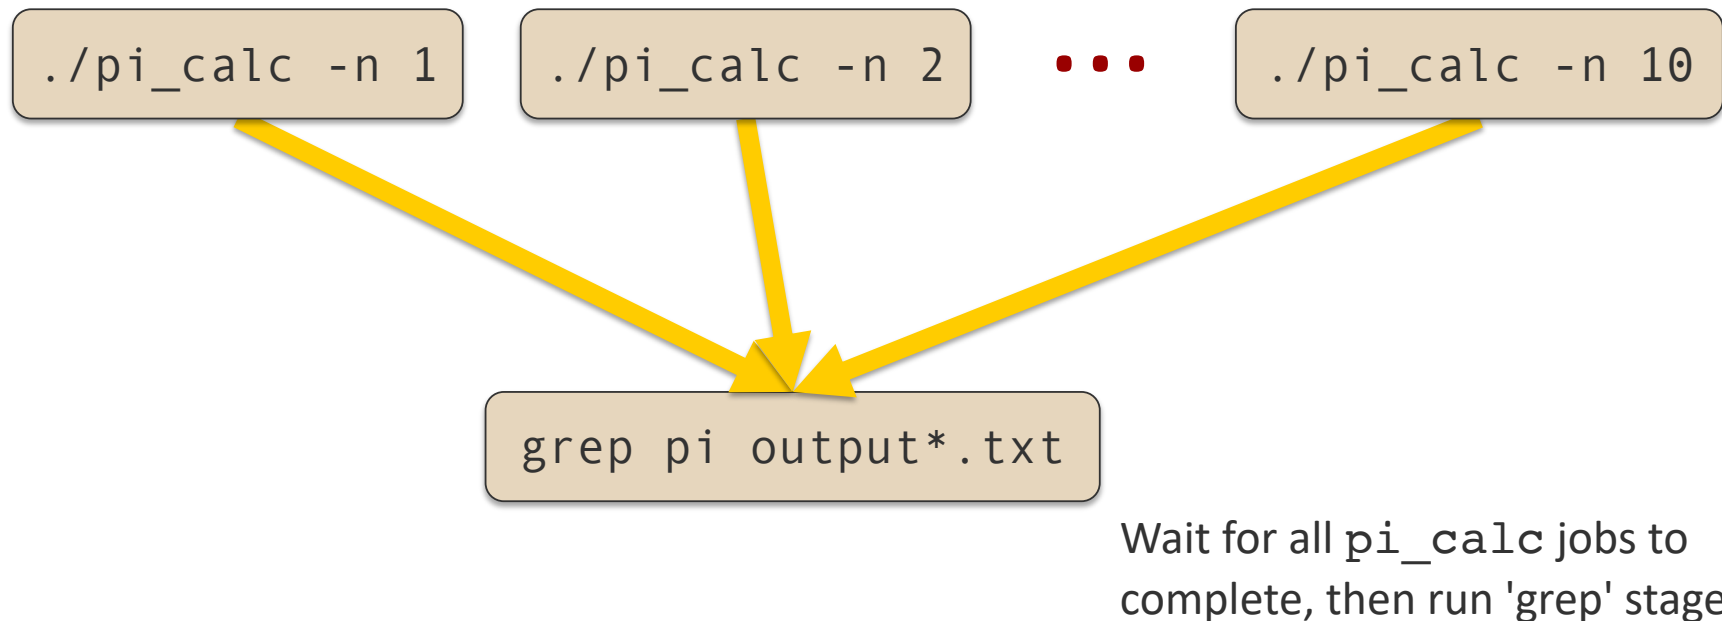
- Much more work per job

■ Cons

- **More complicated**
- Not efficient for heterogeneous run times

Job Dependencies

- You can tell Slurm to put a job on hold until others finish
- Syntax is `#SBATCH --dependency=<dependency_list>`
- Read more on Slurm's [sbatch](#) documentation



Job Dependencies

- Job script

```
#!/bin/bash
#SBATCH --ntasks=1

grep pi output*.txt | sed 's/^.*=//' > summary.txt
```

- Dependency tracker - shell script

```
#!/bin/bash

# Launch first job
slurm_output=$(sbatch job_array.slurm)
# Get job id
dependency=$(echo $slurm_output | awk '{print $NF}')

# Assign dependency
sbatch --depend=afterok:${dependency} summarize.slurm
```

Job Dependencies

- Use dependency.sh to launch ALL jobs in workflow
- [ttroj@hpc-login3 advanced_hpc]\$ myqueue

JOBID	USER	ACCOUNT	PARTITION	NAME	TASKS	CPUS_ PER_ TASK	MIN_ MEMORY	START_ TIME	TIME	TIME_ LIMIT	ST	NODELIST (REASON)
4062911	ttroj	lc_tt1	quick	job_array.slurm	1	1	1G	N/A	0:03	30:00:00	CF	hpc1012
4062912	ttroj	lc_tt1	quick	job_array.slurm	1	1	1G	N/A	0:03	30:00:00	CF	hpc1227
4062913	ttroj	lc_tt1	quick	job_array.slurm	1	1	1G	N/A	0:03	30:00:00	CF	hpc1230
4062905	ttroj	lc_tt1	quick	job_array.slurm	1	1	1G	N/A	0:00	30:00:00	PD	Priority
4062906	ttroj	lc_tt1	quick	summarize.slurm	1	1	1G	N/A	0:00	30:00:00	PD	Dependency
4062909	ttroj	lc_tt1	quick	job_array.slurm	1	1	1G	N/A	0:03	30:00:00	R	hpc0971
4062910	ttroj	lc_tt1	quick	job_array.slurm	1	1	1G	N/A	0:03	30:00:00	R	hpc0972
4062911	ttroj	lc_tt1	quick	job_array.slurm	1	1	1G	N/A	0:03	30:00:00	CF	hpc1012
4062912	ttroj	lc_tt1	quick	job_array.slurm	1	1	1G	N/A	0:03	30:00:00	CF	hpc1227
4062913	ttroj	lc_tt1	quick	job_array.slurm	1	1	1G	N/A	0:03	30:00:00	CF	hpc1230

Job Dependencies

```
[ttroj@hpc-login3 job_dependency]$ cat summary.txt
4.000000000000000000
3.600000000000000001
3.4564102564102557
3.3811764705882354
3.3349261138109898
3.3036297331379298
3.2810484527641703
3.2639884944910889
3.2506461552653931
3.2399259889071588
```



Questions? (HPC@USC.EDU)