

# Parallel MATLAB on HPC

**Erin Shaw and Cesar Sul**

**Advanced Cyberinfrastructure Research and Education Facilitation**

**USC Center for High-Performance Computing**

# Welcome HPC Workshop - Setup

- Sign in at goo.gl/Qan5Kg
- Best to have X display forwarding enabled
  - On Mac, download and install XQuartz from [www.quartz.org](http://www.quartz.org)
  - On PC, enable X forwarding within your ssh client
- Log in to hpc-login2 or hpc-login3 with -X or –Y option
  - Go to your /staging directory and create workshop space

```
$ cd /staging/<project>/<user>
$ mkdir matlab
```

# Welcome HPC Workshop - Setup

- Slides and files used are here:
  - /home/rcf-proj/workshop/handouts/HPCParallelMatlab\*.pdf
  - /home/rcf-proj/workshop/matlab/
- To copy slides to your laptop:
  - a) Use SFTP client like Filezilla
  - b) Open Terminal window, type cd ~', and then:

```
scp <your_NetId>@hpc-login3.usc.edu:/home/rcf-proj/  
workshop/handouts/HPCParallelMATLAB*.pdf <space> .
```

Note: This is all one command and there is a “dot”\* at the end!

\*“dot” is a symbol that resolves to/is replaced by the path to your current directory.

Note: This will download the file to your home directory (on your mac).

# HPC Facilitation

- Request assistance
  - Email [hpc@usc.edu](mailto:hpc@usc.edu) (email again!)
  - Drop-in to office hours
    - UPC: every Tuesday, 2:30p (LVL 3M)
    - HSC: 1<sup>st</sup> & 3<sup>rd</sup> Thursdays (NML 203, SSB 106)
- Learn more!
  - Webpages online at <https://hpcc.usc.edu>
  - Request a consultation ([Getting Help => Office Hours & Consultations](#))
  - Attend a workshop



# Outline

- Introduction
  - Prerequisites, setup workspace and launch MATLAB
  - Introduction to parallel MATLAB and parfor
- Run MATLAB on single node with Parallel Computing Toolbox (PCT)
- Run MATLAB on multiple nodes with Parallel Server
- Run multiple MATLAB jobs with Slurm
- Beyond parfor
  - labIndex, Composite, and Broadcast&Receive

# Outline

- **Introduction**
  - Prerequisites, setup workspace and launch MATLAB
  - Introduction to parallel MATLAB and parfor
- Run MATLAB on single node with Parallel Computing Toolbox (PCT)
- Run MATLAB on multiple nodes with Parallel Server
- Run multiple MATLAB jobs with Slurm
- Beyond parfor
  - labIndex, Composite, and Broadcast&Receive

# Workspace setup

- Start in your /staging directory

```
$ cd /staging/<project>/<user>  
$ mkdir matlab
```

- Make a workshop directory and copy MATLAB examples

```
$ cd matlab  
$ cp -r /home/rcf-proj/workshop/matlab/* .  
$ ls  
helper_script_examples/  standalone_examples/
```

- Make a storage directory for temporary working files

```
$ mkdir matlab_storage  
$ ls
```

# Request an interactive compute node

```
$ salloc --ntasks=1 --cpus-per-task=8 --time=1:00:00 --mem-per-cpu=2gb  
  
--ntasks=1          # keep all cpus on 1 node  
--cpus-per-task=16 # use 16 cpus  
--time=1:00:00      # request up to an hour interactive session  
--mem-per-cpu=2g   # request 2GB of memory per cpu
```



**USC** University of  
Southern California

**USC ITS**  
Information Technology Services

# Wait for compute node

```
[tt@hpc-login3 matlab]$ salloc --ntasks=8 --time=1:00:00 --mem-per-cpu=2g
salloc: Pending job allocation 1084759
salloc: job 1084759 queued and waiting for resources
salloc: job 1084759 has been allocated resources
salloc: Granted job allocation 1084759
salloc: Waiting for resource configuration
salloc: Nodes hpc3779 are ready for job
-----
Begin SLURM Prolog Thu Jun 21 14:52:49 2018
Job ID:          1084759
Username:        tt
Accountname:    lc_tt1
Name:           sh
Partition:      quick
Nodes:          hpc3779
TasksPerNode:   8
CPUSPerTask:   Default[1]
TMPDIR:         /tmp/1084759.quick
SCRATCHDIR:     /staging/scratch/1084759
Cluster:        uschpc
HSDA Account:  false
End SLURM Prolog
-----
[tt@hpc3779 matlab]$ source /usr/usc/matlab/R2018b/setup.sh
[tt@hpc3779 matlab]$ matlab -nodisplay
```

# Outline

- **Introduction**
  - Prerequisites, setup workspace and launch MATLAB
  - **Introduction to parallel MATLAB and parfor**
- Run MATLAB on single node with Parallel Computing Toolbox (PCT)
- Run MATLAB on multiple nodes with Parallel Server
- Run multiple MATLAB jobs with Slurm
- Beyond parfor
  - labIndex, Composite, and Broadcast&Receive

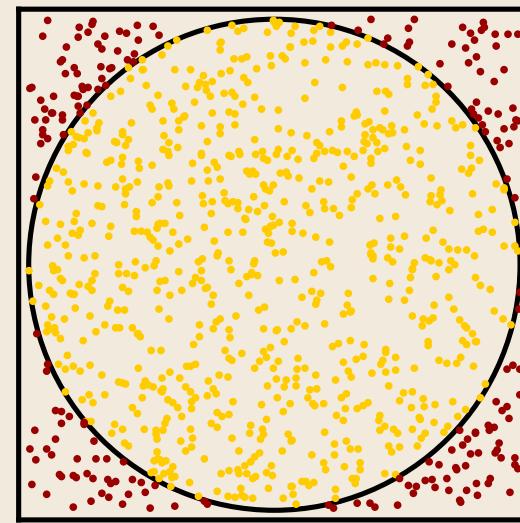
# Parallel MATLAB

- USC has a site-wide license for MATLAB
  - Students can use laptop version free of charge
  - HPC has 1000 licenses for researchers
- HPC license includes use of:
  - MATLAB Parallel Computing Toolbox (PCT) for single-node multi-core jobs
  - MATLAB Parallel Server for multi-node multi-core jobs
- HPC use requires special scripts for Slurm integration
  - MATLAB and Slurm now integrated (>=2018a)
  - /usr/usc/matlab/R2018b/SlurmIntegrationScripts

# Parallel for loop (**parfor**)

```
% EstimatePi.m
% Calculate the value of Pi using a Monte Carlo simulation

max=1e9;
tic;
n=0;
parfor i = 1:max
    x=rand;
    y=rand;
    if (x^2 + y^2 < 1.0)
        n=n+1;
    end
end
elapsedTime = toc;
pi = (4.0 * n / max);
```



# Outline

- Introduction
  - Prerequisites, setup workspace and launch MATLAB
  - Introduction to parallel MATLAB and parfor
- **Run MATLAB on single node with Parallel Computing Toolbox (PCT)**
- Run MATLAB on multiple nodes with Parallel Server
- Run multiple MATLAB jobs with Slurm
- Beyond parfor
  - labIndex, Composite, and Broadcast&Receive

# Run on single node with PCT

- Simplest way to get started is to add these lines to top of your script

```
cluster=parallel.cluster.Local();
cluster.JobStorageLocation='storage_dir';
[a,b]=evalc('system('''nproc --all''')');
cluster.NumWorkers=str2num(a);
```

- **storage\_dir** is absolute path to directory for temporary files

- Start up worker pool with this line in script

```
pool=parpool(cluster,N)
```

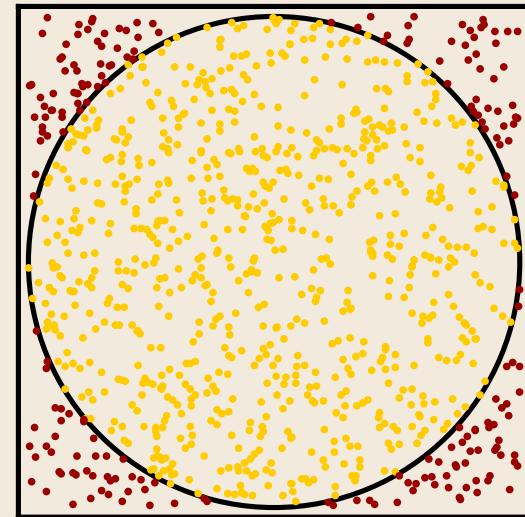
- The value of **N** must be 1 less than total number of cpus allocated

# Create parallel cluster and start parallel pool

```
% EstimatePi.m
% Calculate the value of Pi using a Monte Carlo simulation
cluster=parallel.cluster.Local();
cluster.JobStorageLocation=storage_dir;
[a,b]=evalc('system('''nproc --all''')');
cluster.NumWorkers=str2num(a);

pool=parpool(cluster,N)

max=1e9;
tic;
n=0;
parfor i = 1:max
    x=rand;
    y=rand;
    if (x^2 + y^2 < 1.0)
        n=n+1;
    end
end
elapsedTime = toc;
pi = (4.0 * n / max);
```



# Run EstimatePi.m

- Replace **storage\_dir** in standalone\_examples/estimatePi.m
- Launch MATLAB and run from command line

```
$ matlab -nodisplay -r "estimatePi"  
  
                                < M A T L A B (R) >  
Starting parallel pool (parpool) using the 'local' profile ...  
connected to 16 workers.  
ans =  
Pool with properties:  
    Connected: true  
    NumWorkers: 16  
    Cluster: local  
    AttachedFiles: {}  
    IdleTimeout: 30 minute(s) (30 minutes remaining)  
    SpmdEnabled: true  
ans =  
Elapsed time = 91.211017, pi = 3.1415924292
```



University of  
Southern California

USC ITS

Information Technology Services

# PCT Helper Script

- You can simplify the setup by using helper scripts

```
get_LOCAL_cluster(storage_dir)
```

```
pool=parpool(cluster,N)
```

- **storage\_dir** is absolute path to directory for temporary files
- The value of **N** must be 1 less than total number of cpus allocated

- You must define `get_LOCAL_cluster`

- There is a copy of `get_LOCAL_cluster.m` in `helper_script_examples`
- Copy it to your `~/matlab` or working directory to use

```
cp /home/rcf-proj/workshop/matlab/get_LOCAL_cluster.m  
~/matlab/
```



**USC** University of  
Southern California

**USC ITS**

Information Technology Services

# get\_LOCAL\_cluster.m

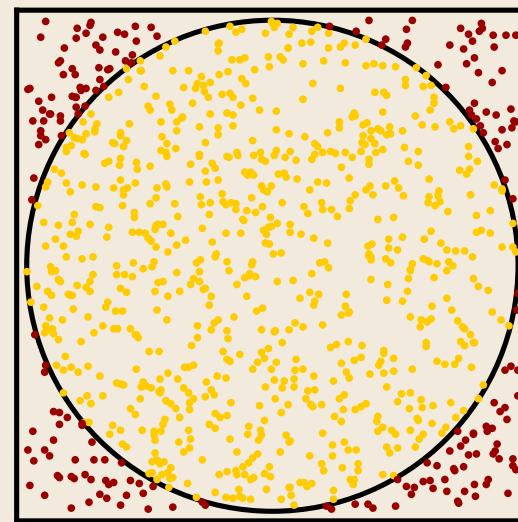
```
% get_LOCAL_cluster(storage_dir)
% HPC helper function to create a single node MATLAB cluster under Slurm.
% Place this file in MATLAB's search path and call it from your program.
% Or place the lines of the function directly within your program.
%
% Arguments:
% storage_dir: The path to an *existing* directory where MATLAB can store
%               files, e.g., "/home/rcf-proj/tt/ttrogan/matlab_storage".
%
function cluster = get_LOCAL_cluster(storage_dir)
    cluster=parallel.cluster.Local();
    cluster.JobStorageLocation=storage_dir;
    [a,b]=evalc('system('''nproc --all''')');
    cluster.NumWorkers=str2num(a);
end
```

# PCT Helper Script: EstimatePi.m

```
% Calculate the value of Pi using a Monte Carlo simulation
```

```
get_LOCAL_cluster(storage_dir)
pool=parpool(cluster,N)

max=1e9;
tic;
n=0;
parfor i = 1:max
    x=rand;
    y=rand;
    if (x^2 + y^2 < 1.0)
        n=n+1;
    end
end
elapsedTime = toc;
pi = (4.0 * n / max);
```



# Submit EstimatePi.m to queue

```
#!/bin/bash

#SBATCH --ntasks=1
#SBATCH --cpus-per-task=16
#SBATCH --mem-per-cpu=1g
#SBATCH --time=00:20:00
#SBATCH --export=none

source /usr/usc/matlab/R2018a/setup.sh

#Suppress time warning
export TZ=America/Los_Angeles

matlab -nodisplay -r "estimatePi"
```



**USC** University of  
Southern California

**USC ITS**  
Information Technology Services

# Outline

- Introduction
  - Prerequisites, setup workspace and launch MATLAB
  - Introduction to parallel MATLAB and parfor
- Run MATLAB on single node with Parallel Computing Toolbox (PCT)
- **Run MATLAB on multiple nodes with Parallel Server**
- Run multiple MATLAB jobs with Slurm
- Beyond parfor
  - labIndex, Composite, and Broadcast&Receive

# Run on multiple nodes with Parallel Server

- R2018/home/usc/matlab/R2018b/SlurmIntegrationScripts
- Simplest way to get started, add this to top of Matlab script

```
cluster = parallel.cluster.Generic;
set(cluster,'JobStorageLocation', 'storage_dir');
set(cluster,'HasSharedFilesystem', true);
set(cluster,'IntegrationScriptsLocation','scripts_dir');
cluster.AdditionalProperties.SlurmArgs='sbatch_args';
```

|                    |  |
|--------------------|--|
| <b>storage_dir</b> | Path to where temp Matlab files are saved  |
| <b>scripts_dir</b> | Path to Slurm integration scripts:<br>/usr/usc/matlab/ <b>version</b> /SlurmIntegrationScr<br>ipts |
| <b>sbatch_args</b> | Arguments for job scheduler  |

# Run on multiple nodes

- Matlab will launch a **second** job after the parpool command

```
pool=parpool(cluster,N)
```

- Memory requests, wallclock time specified in **sbatch\_args**
- Matlab will use **N** to determine number of cpus to request.



**USC** University of  
Southern California

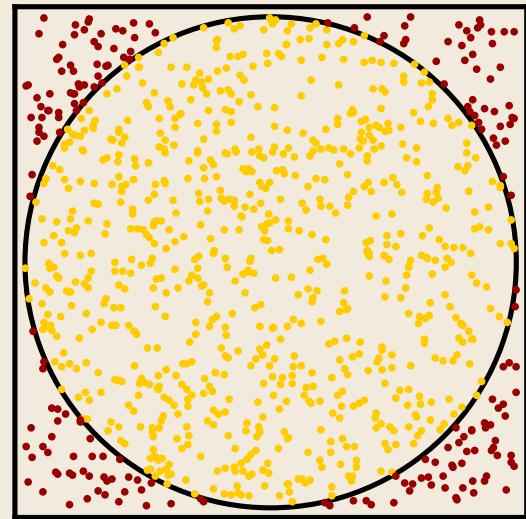
**USC ITS**  
Information Technology Services

# Setup cluster and start parallel pool

```
% Calculate the value of Pi using a Monte Carlo simulation
cluster = parallel.cluster.Generic;
set(cluster, 'JobStorageLocation', storage_dir);
set(cluster, 'HasSharedFilesystem', true);
set(cluster, 'IntegrationScriptsLocation', scripts_dir);
cluster.AdditionalProperties.SlurmArgs=sbatch_args;

pool=parpool(cluster,N)

max=1e9;
tic;
n=0;
parfor i = 1:max
    x=rand;
    y=rand;
    if (x^2 + y^2 < 1.0)
        n=n+1;
    end
end
elapsedTime = toc;
pi = (4.0 * n / max);
```



# Run EstimatePi.m

- Launch Matlab and run from command line

```
$ matlab -nodisplay -r "estimatePi"
```

```
< M A T L A B (R) >
```

```
Warning: The IntegrationsScriptsLocation  
/usr/usc/matlab/R2018a/SlurmIntegrationScripts has been put at  
the top of the path.  
connected to 7 workers.
```

```
ans =
```

```
Pool with properties:
```

```
    Connected: true  
    NumWorkers: 7  
    Cluster: Generic  
    AttachedFiles: {}  
    AutoAddClientPath: true  
    IdleTimeout: 30 minutes (30 minutes remaining)  
    SpmdEnabled: true
```

```
EnvironmentVariables: {}
```

```
pi = 3.1416, elapsedTime = 2.5654
```



# Parallel Server Helper Script

- If you want your scripts to be simple

```
get_SLURM_cluster(storage_dir,scripts_dir,sbatch_args)  
pool=parpool(cluster,N)
```

|                    |  |
|--------------------|--|
| <b>storage_dir</b> | Path to where temp Matlab files are saved  |
| <b>scripts_dir</b> | Path to Slurm integration scripts:<br><code>/usr/usc/matlab/version/SlurmIntegrationScripts</code> |
| <b>sbatch_args</b> | Arguments for job scheduler  |
| <b>N</b>           | Number of workers to use   |

- Copy `get_SLURM_cluster.m` to your home directory

```
cp /home/rcf-proj/workshop/matlab/get_SLURM_cluster.m ~/matlab/
```

# get\_SLURM\_cluster.m

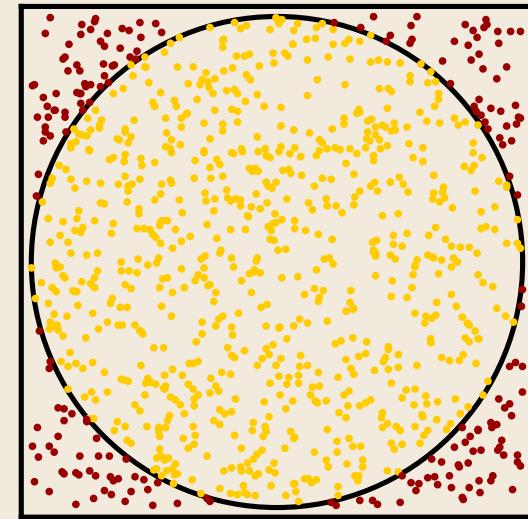
```
%%%%%%
% get_SLURM_cluster(storage_dir,scripts_dir,sbatch_args)
% HPC helper function to create a multi-node MATLAB cluster under Slurm.
% Place this file in MATLAB's search path and call it from your program.
% Or place the lines of the function directly within your program.
%
% Arguments:
%   storage_dir: The path to an *existing* directory where MATLAB can store
%                 files, e.g., "/home/rcf-proj/tt/trojan/matlab_storage".
%   scripts_dir: The path to the /usr/usc/matlab//SlurmIntegrationScripts
%                 directory, e.g.,
"/usr/usc/matlab/R2018a/SlurmIntegrationScripts".
%   sbatch_arg: A string to pass to slurm with all slurm options *except*
%               ntasks, e.g., '--time=12:00:00 --partition=scec --mem-per-
cpu=2G'.
%%%%%
function cluster = get_SLURM_cluster(storage_dir,scripts_dir,sbatch_args)
    cluster = parallel.cluster.Generic;
    set(cluster,'JobStorageLocation', storage_dir);
    set(cluster,'HasSharedFilesystem', true);
    set(cluster,'IntegrationScriptsLocation',scripts_dir);
    cluster.AdditionalProperties.SlurmArgs=sbatch_args;
end
```

# Helper Script: estimatePi.m

```
% Calculate the value of Pi using a Monte Carlo simulation
cluster = get_SLURM_cluster('storage_dir', ...
'/usr/usc/matlab/R2018a/SlurmIntegrationScripts', ...
'--mem-per-cpu=1g --time=00:30:00');

pool=parpool(cluster,N)

max=1e9;
tic;
n=0;
parfor i = 1:max
    x=rand;
    y=rand;
    if (x^2 + y^2 < 1.0)
        n=n+1;
    end
end
elapsedTime = toc;
pi = (4.0 * n / max);
```



# Submit EstimatePi.m to queue

```
#!/bin/bash

#SBATCH --ntasks=1
#SBATCH --cpus-per-task=8
#SBATCH --mem-per-cpu=1g
#SBATCH --time=00:20:00
#SBATCH --export=none

source /usr/usc/matlab/R2018a/setup.sh

#Suppress time warning
export TZ=America/Los_Angeles

matlab -nodisplay -r "estimatePi"
```



**USC** University of  
Southern California

**USC ITS**  
Information Technology Services

# Outline

- Introduction
  - Prerequisites, setup workspace and launch MATLAB
  - Introduction to parallel MATLAB and parfor
- Run MATLAB on single node with Parallel Computing Toolbox (PCT)
- Run MATLAB on multiple nodes with Parallel Server
- **Run multiple MATLAB jobs with Slurm**
- Beyond parfor
  - labIndex, Composite, and Broadcast&Receive

# Slurm Job Array

- If one node of your parallel cluster fails, the whole cluster fails.
  - This can be a problem.
  - If your workers are running independent calculations, you can use Slurm's job array construct instead
  - To run the following 100 times: `sbatch --array=1-100 myjob.sl  
$ matlab –nodisplay –r “mymatlab”`
- Partitions have per user limits on cpu, job and mem allocations
  - Main: Per user: up to 200 cpus, 10 jobs
  - Scavenge: Per user: up to 500 cpus, 2500G RAM (no limit on jobs)
- When limits are reached your job(s) go to the back of the queue
  - They are still in the queue!
  - You do not have to manage them, they will eventually run

# Slurm Job Array

```
$ cat matlab.sl
```

```
#!/bin/bash
#SBATCH --ntasks=100
cd $SLURM_SUBMIT_DIR
source /usr/usc/matlab/R2018b/setup.sh
matlab –nodisplay –r “disp(‘hello’);exit”
```

```
$ sbatch --array=1-10 --partition=scavenge matlab.sl
```

- What happens?
  - Hits max CPUs after 5 jobs
- Now set --ntasks=10 and –array=1-100, what happens?
  - Would hit max CPUs after 50 jobs. Instead hits max jobs at 10 jobs

# Slurm Job Array

```
$ sbatch --array=1-100 --partition=scavenge matlab.sl      #ntasks=100
```

```
$ squeue -u $USER
```

| JOBid           | Partition | Name  | User    | St | Time | Nodes | NodeList(Reason) |
|-----------------|-----------|-------|---------|----|------|-------|------------------|
| 3724946_[1-100] | scavenge  | array | ttrojan | PD | 0:00 | 13    | (Priority)       |

```
$ squeue -u $USER
```

| JOBid           | Partition | Name  | User     | St | Time | Nodes | NodeList(Reason) |
|-----------------|-----------|-------|----------|----|------|-------|------------------|
| 3724946_[6-100] | scavenge  | array | erinshaw | PD | 0:00 | 13    | (Priority)       |



USC University of  
Southern California

USC ITS  
Information Technology Services

# Outline

- Introduction
  - Prerequisites, setup workspace and launch MATLAB
  - Introduction to parallel MATLAB and parfor
- Run MATLAB on single node with Parallel Computing Toolbox (PCT)
- Run MATLAB on multiple nodes with Parallel Server
- Run multiple MATLAB jobs with Slurm
- **Beyond parfor**
  - labIndex, Composite, and Broadcast&Receive

# Beyond Parfor

- Parfor is the most basic way that Matlab can parallelize your code
- Sometimes you'll need to have more control over what gets run on each worker
- Here are some examples to give you an idea of the capabilities of the Parallel MATLAB
- Files saved in /home/rcf-proj/workshop/matlab/



**USC** University of  
Southern California

**USC ITS**  
Information Technology Services

# submitTasks.m

```
% Demo create job and task and submit  
  
% Create cluster object (single or multi node)  
cluster=get_LOCAL_cluster(storage_dir)  
  
% create cluster with 2 workers  
cluster.parpool(2)  
  
% create job  
j = c.createJob;  
  
a = rand(4);  
% create task  
j.createTask(@sin,1,{a});  
j.createTask(@cos,1,{a});  
  
% have two independent jobs to submit  
j.submit;  
  
%wait till job finish  
j.wait('finished');  
pause(20);  
  
% fetch output  
result = j.fetchOutputs;  
f = [result{:}];  
  
% save data  
save('submitTasks.out','f', '-ascii' );  
delete(gcp('nocreate'))
```

# labIndex.m

```
% Demo of parpool, Composite and labindex  
  
% Create cluster object (single or multi node)  
cluster=get_LOCAL_cluster(storage_dir)  
  
% request 3 workers  
cluster.parpool(3)
```

```
% compute process  
a =Composite();  
spmd  
    temp=labIndex*ones(10);  
    a = temp*temp;  
end  
  
% save worker results  
d=[a{:}];  
Save('labIndex.out', 'd', '-ascii' );  
  
%cleanup  
delete(gcp('nocreate'))
```

# broadcastReceive.m

```
% Demo labindex and mpi communication
% Create cluster object (single or multi node)
cluster=get_LOCAL_cluster(storage_dir)

% request workers
cluster.parpool(3)

disp('Test one')
spmd
switch labindex
    case 1
        A = 1:5;
        labSend(A,2,199); %send to worker 2
    case 2
        A = labReceive(1,199); %receive from worker 1
    otherwise
        A = 11:15; %create his own
    end
end
A{[:]}

%what does this do?
disp('Test two')

spmd
if labindex == 1
    A=labBroadcast (1,1:5); %broadcast to all
else
    A=labBroadcast (1); %receive broadcast
end
end
A{[:]}
delete(gcp('nocreate')) )
```

# The End

## Questions?

- To view licenses available
  - Type on command line

```
$ /auto/usc/matlab/<version>/etc/glnxa64/lmutil lmstat -a -c 27000@hpc-licenses.usc.edu | grep MATLAB
```

Users of MATLAB: (Total of 14146 licenses issued; Total of 4 licenses in use)

Users of MATLAB\_Distrib\_Comp\_Engine: (Total of 1000 licenses issued; Total of 0 licenses in use)



**USC** University of  
Southern California

**USC ITS**  
Information Technology Services